

# Iterative LQR & Model Predictive Control


Sanjiban Choudhury

TAs: Matthew Rockett, Gilwoo Lee, Matt Schmittle

# Table of Controllers

|              | Control Law | Uses model    | Stability Guarantee    | Minimize Cost |
|--------------|-------------|---------------|------------------------|---------------|
| PID          |             | No            | No                     | No            |
| Pure Pursuit |             | Circular arcs | Yes - with assumptions | No            |
| Lyapunov     |             | Non-linear    | Yes                    | No            |
| LQR          |             | Linear        | Yes                    | Quadratic     |


# Can we use LQR to swing up a pendulum?



# Can we use LQR to swing up a pendulum?

No!

(Large angles imply  
large linearization error)




# Can we use LQR to swing up a pendulum?

No!

(Large angles imply  
large linearization error)

But we can track  
a reference swing up trajectory  
(small linearization error)



But, first we need to talk  
about time-varying systems

# Today's objectives

1. LQR for time-varying systems
2. Trajectory following with iLQR
3. General nonlinear trajectory optimization with iLQR
4. Model predictive control (MPC)

# LQR for Time-Varying Dynamical Systems

$$x_{t+1} = A_t x_t + B_t u_t$$



# LQR for Time-Varying Dynamical Systems

$$x_{t+1} = A_t x_t + B_t u_t$$

$$c(x_t, u_t) = x_t^T Q_t x_t + u_t^T R_t u_t$$

# LQR for Time-Varying Dynamical Systems

$$x_{t+1} = A_t x_t + B_t u_t$$

$$c(x_t, u_t) = x_t^T Q_t x_t + u_t^T R_t u_t$$

Straight forward to get LQR equations


$$K_t = -(R_t + B_t^T V_{t+1} B_t)^{-1} B_t^T V_{t+1} A_t$$

$$V_t = Q_t + K_t^T R_t K_t + (A_t + B_t K_t)^T V_{t+1} (A_t + B_t K_t)$$


Why do we care about time-varying?

Ans: Linearization about a trajectory

# Trajectory tracking for stationary rolls?

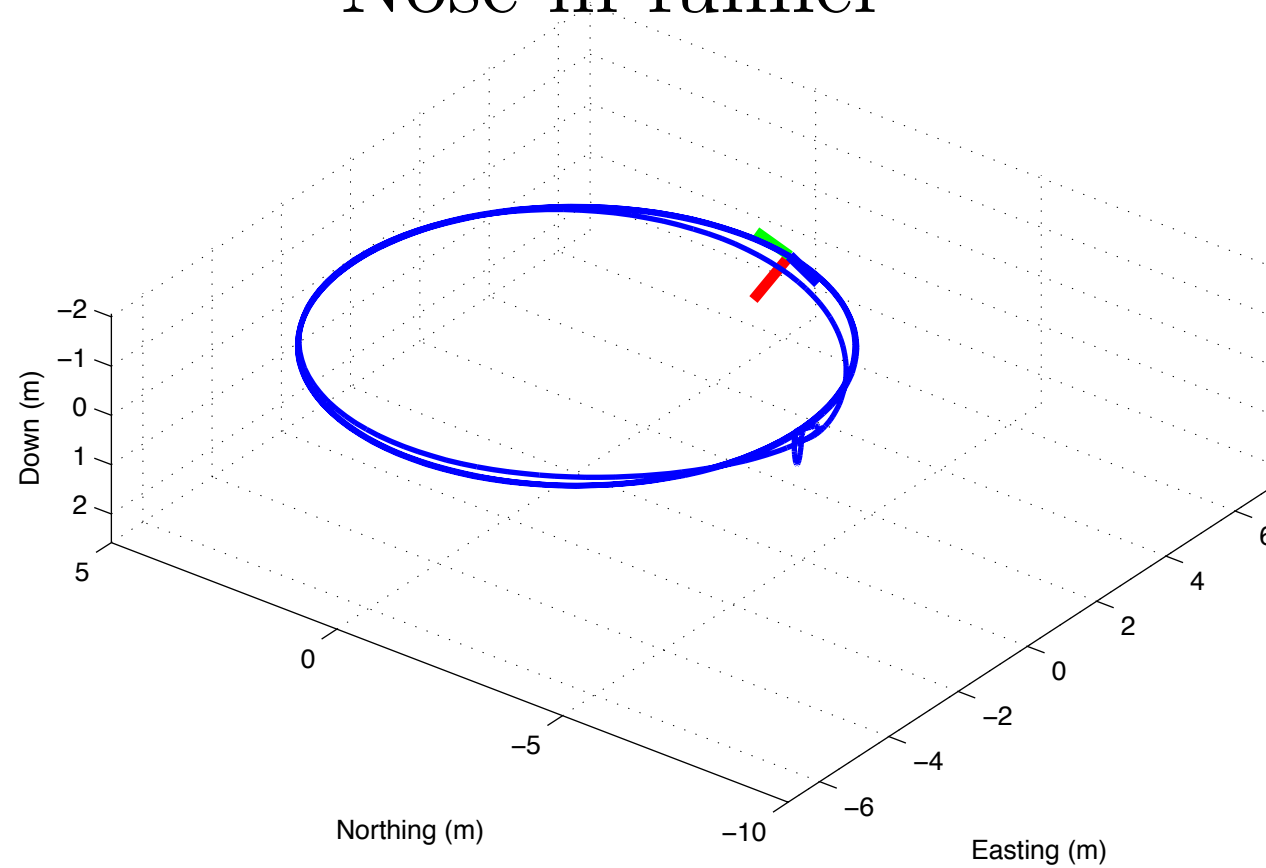


# Trajectory tracking for stationary rolls?




# How do we get such behaviors?

Nose-in funnel



Stationary rolls



# Task: Minimize tracking error

$$\min_{u_0, u_1, \dots, u_{T-1}} \sum_{t=0}^{T-1} c(x_t, u_t)$$


$$\text{subject to } x_{t+1} = f(x_t, u_t) \quad \forall t$$

In this scenario, cost is simply a quadratic tracking cost


Why is this a hard optimization problem?

# Iterative LQR (iLQR)


Start by guessing a control sequence, Forward simulate dynamics,  
**Linearize** about trajectory, Solve for new control sequence  
and repeat!



$i=0$



$i=10$




$i=100$



# Step 1: Get a reference trajectory

$$x_0^{ref}, u_0^{ref}, x_1^{ref}, u_1^{ref}, \dots, x_{T-1}^{ref}, u_{T-1}^{ref}$$



**Note:** Simply executing open loop trajectory wont work!

# Step 2: Initialize your algorithm

Choose initial trajectory at iteration 0 to linearize about

$$x^0(t), u^0(t) = \{x_0^0, u_0^0, x_1^0, u_1^0, \dots, x_{T-1}^0, u_{T-1}^0\}$$

# Step 2: Initialize your algorithm

Choose initial trajectory at iteration 0 to linearize about

$$x^0(t), u^0(t) = \{x_0^0, u_0^0, x_1^0, u_1^0, \dots, x_{T-1}^0, u_{T-1}^0\}$$

It's a good idea to choose the reference trajectory

# Step 2: Initialize your algorithm

Choose initial trajectory at iteration 0 to linearize about

$$x^0(t), u^0(t) = \{x_0^0, u_0^0, x_1^0, u_1^0, \dots, x_{T-1}^0, u_{T-1}^0\}$$

It's a good idea to choose the reference trajectory

Initialization is very important!

We will be perturbing this initial trajectory

# Step 3: Linearize your dynamics!

At a given iteration  $i$ , we are going to linearize about

$$x_0^i, u_0^i, x_1^i, \dots$$

# Step 3: Linearize your dynamics!

At a given iteration  $i$ , we are going to linearize about


$$x_0^i, u_0^i, x_1^i, \dots$$

Change of variable - we will track the delta perturbations

$$\delta x_t = x_t - x_t^i$$

$$\delta u_t = u_t - u_t^i$$

# Step 3: Linearize your dynamics!



# Step 3: Linearize your dynamics!

$$\delta x_t = x_t - x_t^i$$

$$\delta u_t = u_t - u_t^i$$



# Step 3: Linearize your dynamics!

$$\delta x_t = x_t - x_t^i$$

$$\delta u_t = u_t - u_t^i$$

$$\delta x_{t+1} = A_t \delta x_t + B_t \delta u_t + (f(x_t^i, u_t^i) - x_{t+1}^i)$$

$$A_t = \left. \frac{\partial f}{\partial x} \right|_{x_t^i}$$

$$B_t = \left. \frac{\partial f}{\partial u} \right|_{u_t^i}$$

# Step 3: Linearize your dynamics!

$$\delta x_t = x_t - x_t^i$$

$$\delta u_t = u_t - u_t^i$$

This is an affine system, not linear

$$\delta x_{t+1} = A_t \delta x_t + B_t \delta u_t + (f(x_t^i, u_t^i) - x_{t+1}^i)$$

$$A_t = \left. \frac{\partial f}{\partial x} \right|_{x_t^i}$$

$$B_t = \left. \frac{\partial f}{\partial u} \right|_{u_t^i}$$

# Step 3: Linearize your dynamics!

Homogenous coordinate system  $\begin{bmatrix} \delta x_t \\ 1 \end{bmatrix}$

Affine dynamics is now linear!

$$\begin{bmatrix} \delta x_{t+1} \\ 1 \end{bmatrix} = \underbrace{\begin{bmatrix} A_{t+1} & f(x_t^i, u_t^i) - x_{t+1}^i \\ 0 & 1 \end{bmatrix}}_{\tilde{A}_t} \begin{bmatrix} \delta x_t \\ 1 \end{bmatrix} + \underbrace{\begin{bmatrix} B_{t+1} \\ 0 \end{bmatrix}}_{\tilde{B}_t} \delta u_t$$

# Step 4: Quadraticize cost about trajectory

Our cost function is already quadratic, otherwise we would apply  
Taylor expansion

$$c(x_t, u_t) = (x_t - x_t^{ref})^T Q (x_t - x_t^{ref}) + (u_t - u_t^{ref})^T R (u_t - u_t^{ref})$$

# Step 4: Quadraticize cost about trajectory

Our cost function is already quadratic, otherwise we would apply  
Taylor expansion


$$\begin{aligned} c(x_t, u_t) &= (x_t - x_t^{ref})^T Q (x_t - x_t^{ref}) + (u_t - u_t^{ref})^T R (u_t - u_t^{ref}) \\ &= \begin{bmatrix} \delta x_t \\ 1 \end{bmatrix}^T \begin{bmatrix} Q & Q(x_t^i - x_t^{ref}) \\ (x_t^i - x_t^{ref})^T Q & (x_t^i - x_t^{ref})^T (x_t^i - x_t^{ref}) \end{bmatrix} \begin{bmatrix} \delta x_t \\ 1 \end{bmatrix} \\ &\quad \tilde{Q}_t \\ &\quad + \\ &\quad \begin{bmatrix} \delta u_t \\ 1 \end{bmatrix}^T \begin{bmatrix} R & R(u_t^i - u_t^{ref}) \\ (u_t^i - u_t^{ref})^T R & (u_t^i - u_t^{ref})^T (u_t^i - u_t^{ref}) \end{bmatrix} \begin{bmatrix} \delta u_t \\ 1 \end{bmatrix} \\ &\quad \tilde{R}_t \end{aligned}$$

# We have all the ingredients to call LQR!

$$\tilde{K}_t = -(\tilde{R}_t + \tilde{B}_t^T \tilde{V}_{t+1} \tilde{B}_t)^{-1} \tilde{B}_t^T \tilde{V}_{t+1} \tilde{A}_t$$


similarly calculate the value function ...

# Step 5: Do a backward pass



Calculate controller gains for all time steps

**Step 6:** Do a forward pass to get new trajectory



Compute  
control action


$$u_t^{i+1} = u_t^i + \tilde{K}_t \begin{bmatrix} x_t^{i+1} - x_t^i \\ 1 \end{bmatrix}$$

Apply dynamics

$$x_t^{i+1} = f(x_t^{i+1}, u_t^{i+1})$$



# Step 6: Do a forward pass to get new trajectory




Compute  
control action

$$u_t^{i+1} = u_t^i + \tilde{K}_t \begin{bmatrix} x_t^{i+1} - x_t^i \\ 1 \end{bmatrix}$$

Apply dynamics

$$x_t^{i+1} = f(x_t^{i+1}, u_t^{i+1})$$

# Step 6: Do a forward pass to get new trajectory




Compute  
control action

$$u_t^{i+1} = u_t^i + \tilde{K}_t \begin{bmatrix} x_t^{i+1} - x_t^i \\ 1 \end{bmatrix}$$

Apply dynamics

$$x_t^{i+1} = f(x_t^{i+1}, u_t^{i+1})$$

# Step 6: Do a forward pass to get new trajectory




Compute  
control action

$$u_t^{i+1} = u_t^i + \tilde{K}_t \begin{bmatrix} x_t^{i+1} - x_t^i \\ 1 \end{bmatrix}$$

Apply dynamics

$$x_t^{i+1} = f(x_t^{i+1}, u_t^{i+1})$$

**Step 6:** Do a forward pass to get new trajectory



Compute  
control action


$$u_t^{i+1} = u_t^i + \tilde{K}_t \begin{bmatrix} x_t^{i+1} - x_t^i \\ 1 \end{bmatrix}$$

Apply dynamics

$$x_t^{i+1} = f(x_t^{i+1}, u_t^{i+1})$$


# Problem: Forward pass will go bonkers

## Why?



# Problem: Forward pass will go bonkers


Why?



Linearization error gets bigger and bigger and bigger

# Problem: Forward pass will go bonkers

Why?



Linearization error gets bigger and bigger and bigger

Remedies: Change cost function to **penalize deviation** from linearization

# Questions



# Questions

1. Can we solve LQR for **continuous** time dynamics?

# Questions

1. Can we solve LQR for **continuous** time dynamics?

Yes! Refer to Continuous Algebraic Ricatti Equations (CARE)

# Questions

1. Can we solve LQR for **continuous** time dynamics?

Yes! Refer to Continuous Algebraic Ricatti Equations (CARE)

2. Can LQR handle **arbitrary costs** (not just tracking)?

# Questions

1. Can we solve LQR for **continuous** time dynamics?

Yes! Refer to Continuous Algebraic Ricatti Equations (CARE)

2. Can LQR handle **arbitrary costs** (not just tracking)?

Yes! Just quadricize the cost

# Questions

1. Can we solve LQR for **continuous** time dynamics?

Yes! Refer to Continuous Algebraic Ricatti Equations (CARE)

2. Can LQR handle **arbitrary costs** (not just tracking)?

Yes! Just quadricize the cost

3. What if I want to penalize control **derivatives**?

# Questions

1. Can we solve LQR for **continuous** time dynamics?

Yes! Refer to Continuous Algebraic Ricatti Equations (CARE)

2. Can LQR handle **arbitrary costs** (not just tracking)?

Yes! Just quadricize the cost

3. What if I want to penalize control **derivatives**?

No problem! Add control as part of state space

# Questions

1. Can we solve LQR for **continuous** time dynamics?

Yes! Refer to Continuous Algebraic Ricatti Equations (CARE)

2. Can LQR handle **arbitrary costs** (not just tracking)?

Yes! Just quadricize the cost

3. What if I want to penalize control **derivatives**?

No problem! Add control as part of state space

4. Can we handle **noisy** dynamics?

# Questions

1. Can we solve LQR for **continuous** time dynamics?

Yes! Refer to Continuous Algebraic Ricatti Equations (CARE)

2. Can LQR handle **arbitrary costs** (not just tracking)?

Yes! Just quadricize the cost

3. What if I want to penalize control **derivatives**?

No problem! Add control as part of state space

4. Can we handle **noisy** dynamics?

Yes! Gaussian noise does not change the answer



# Table of Controllers

| Control Law  | Uses model    | Stability Guarantee    | Minimize Cost |
|--------------|---------------|------------------------|---------------|
| PID          | No            | No                     | No            |
| Pure Pursuit | Circular arcs | Yes - with assumptions | No            |
| Lyapunov     | Non-linear    | Yes                    | No            |
| LQR          | Linear        | Yes                    | Quadratic     |
| iLQR         | Non-linear    | Yes                    | Yes           |

iLQR is just one technique

It's far from perfect - can't deal with  
model errors / constraints ...

## Model Predictive Control (MPC)

iLQR

Shooting


MPPI

TrajLib

DDP


CMAES

# Recap: Feedback control framework




Look at current state error and compute control actions

# Recap: Feedback control framework




Look at current state error and compute control actions

# Recap: Feedback control framework



Look at current state error and compute control actions

# Recap: Feedback control framework



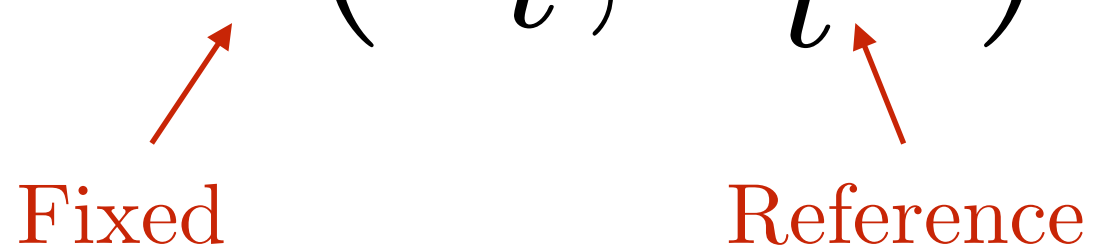
Look at current state error and compute control actions

**Goal:** To drive error to 0 ... to optimally drive it to 0

# Limitations of this framework

A **fixed** control law that looks at **instantaneous** feedback

$$u_t = \pi(x_t, x_t^{\text{ref}})$$

  
Fixed                      Reference

Why is it so difficult to create a magic control law?

# Problem 1: What if we have constraints?

Simple scenario: Car tracking a straight line





# Problem 1: What if we have constraints?



Simple scenario: Car tracking a straight line



Small error,  
control within  
steering constraints


# Problem 1: What if we have constraints?

Simple scenario: Car tracking a straight line



# Problem 1: What if we have constraints?

Simple scenario: Car tracking a straight line



We could “clamp control command” ...  
but what are the implications?

# General problem: Complex models

Dynamics  $x_{t+1} = f(x_t, u_t)$


Constraints  $g(x_t, u_t) \leq 0$

Such complex models imply we need to:

1. Predict the implications of control actions
2. Do corrections NOW that would affect the future
3. It may not be possible to find one law - might need to predict


# Example: Rough terrain mobility

2560, 2.5 second trajectories sampled  
with cost-weighted average @ 60 Hz




# Example: Rough terrain mobility

2560, 2.5 second trajectories sampled  
with cost-weighted average @ 60 Hz



## Problem 2: What if some errors are worse than others?



We need a cost function that penalizes states non-uniformly

Key Idea:

Frame control as an **optimization** problem



# Model predictive control (MPC)

1. Plan a sequence of control actions
2. Predict the set of next states unto a horizon  $H$
3. Evaluate the cost / constraint of the states and controls
4. Optimize the cost

# Model **predictive** control (MPC)

$$\min_{\substack{u_{t+1}, \dots, u_{t+H} \\ \text{(plan till horizon H)}}} \sum_{k=t}^{t+H-1} J(x_k, u_{k+1})$$

(Cost)

$$x_{k+1} = f(x_k, u_{k+1})$$

(**Predict** next state  
with dynamics)

$$g(x_k, u_{k+1}) \leq 0$$


(Constraints)

# Model predictive control (MPC)


$$\min_{u_{t+1}, \dots, u_{t+H}} \sum_{k=t}^{t+H-1} J(x_k, u_{k+1})$$

$$x_{k+1} = f(x_k, u_{k+1})$$

$$g(x_k, u_{k+1}) \leq 0$$




# How are the controls executed?



Step 1: Solve optimization problem to a horizon


# How are the controls executed?



Step 1: Solve optimization problem to a horizon

Step 2: Execute the first control

# How are the controls executed?

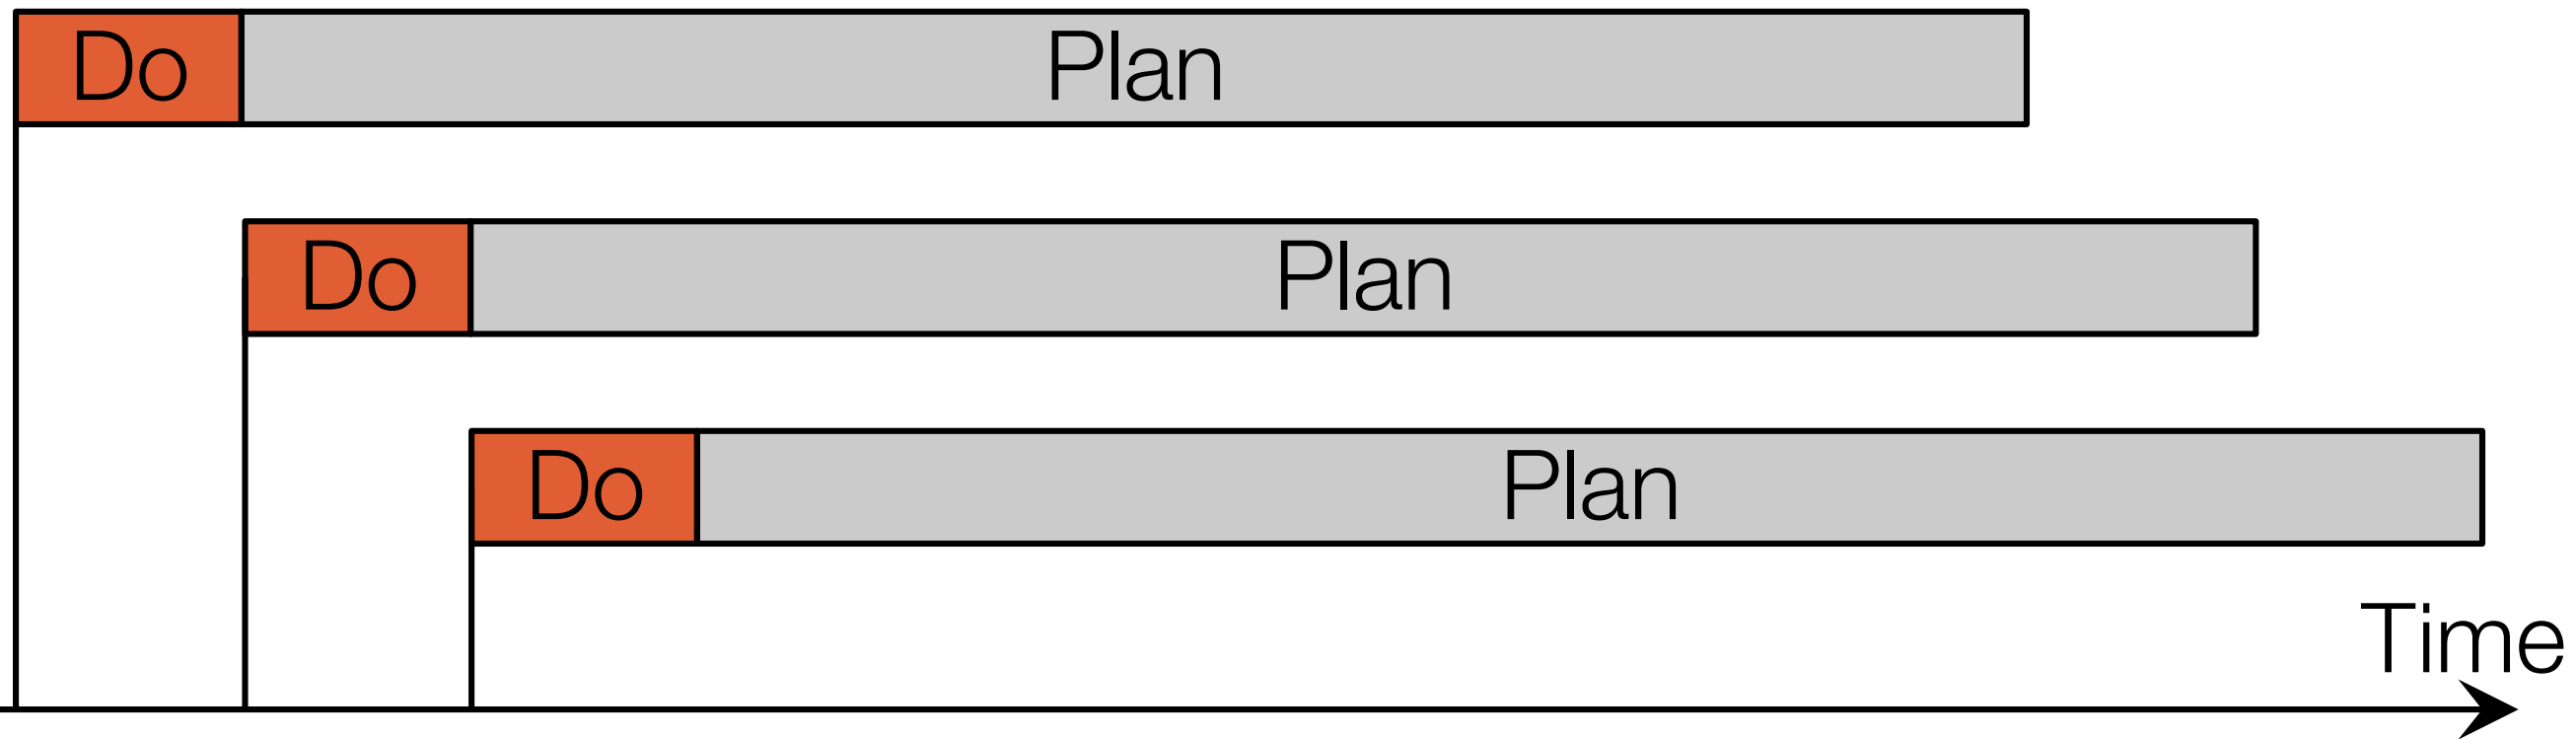


Step 1: Solve optimization problem to a horizon

Step 2: Execute the first control

Step 3: Repeat!

# MPC is a framework




Step 1: Solve optimization problem to a horizon

Step 2: Execute the first control

Step 3: Repeat!


# Why do we need to replan?



What happens if the controls are planned once and executed?



# Why do we need to replan?



What happens if the controls are planned once and executed?

# Model predictive control (MPC)

$$\min_{\substack{u_{t+1}, \dots, u_{t+H} \\ \text{(plan till horizon H)}}} \sum_{k=t}^{t+H-1} J(x_k, u_{k+1})$$

(Cost)

$$x_{k+1} = f(x_k, u_{k+1})$$

(**Predict** next state  
with dynamics)

$$g(x_k, u_{k+1}) \leq 0$$

(Constraints)